


Lecture 6,7 : Thread

From Processes to Threads





The Soul of a Process

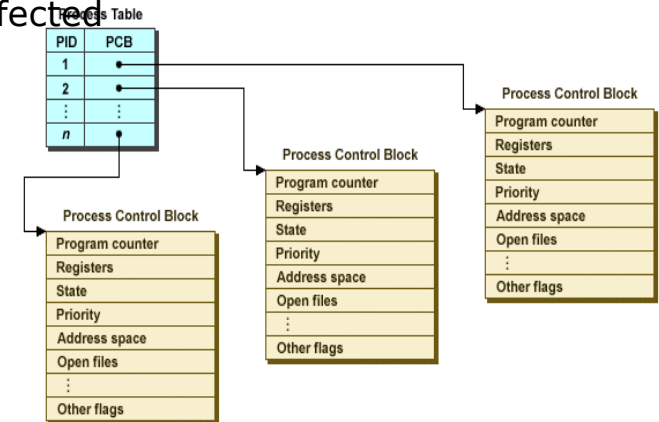
Shared data, has IPC, execution been affected

What is similar in cooperating processes?

They all share the same code and data (address space)

They all share the same privileges

They all share the same resources (files, sockets, etc.)



What don't they share?


Each has its own execution state: PC, SP, and registers

Key idea: Why don't we separate the concept of a process from its execution state?

Process: address space, privileges, resources, etc.

Execution state: PC, SP, registers

Exec state also called **thread of control**, or **thread**





Processes and Threads

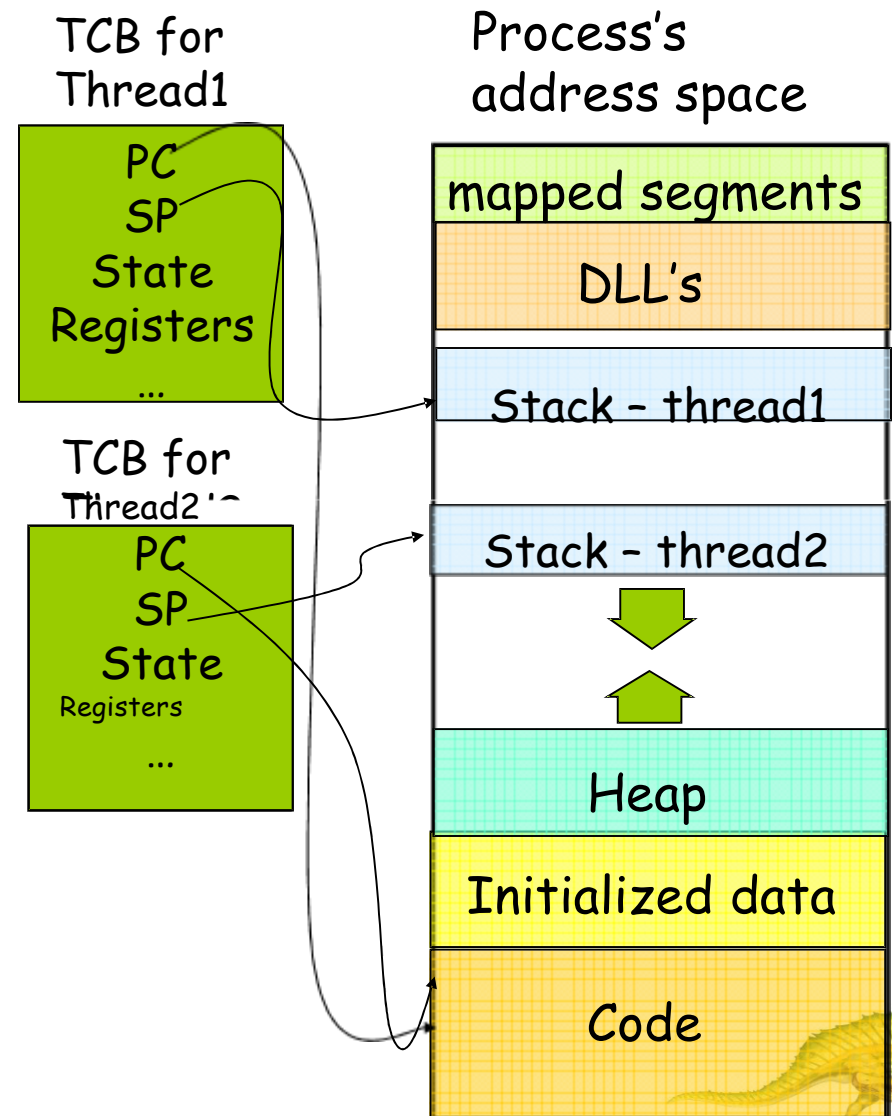
Processes define an address space;
threads share the address space

Process Control Block (PCB)
contains process-specific information

Owner, PID, heap pointer, priority,
active thread, and pointers to thread
information

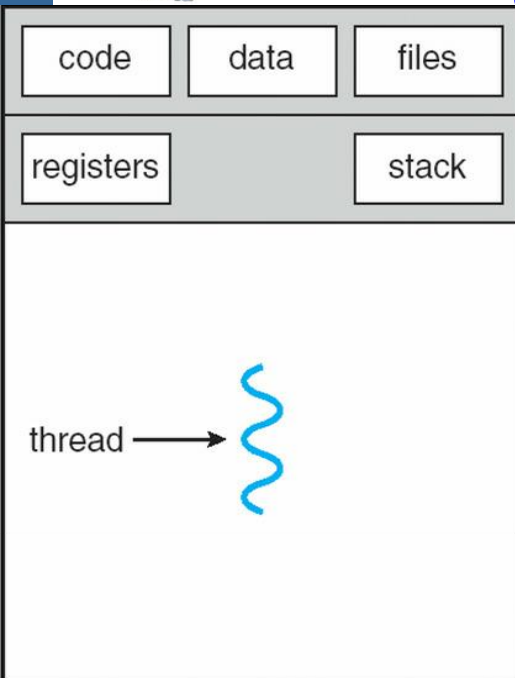
Thread Control Block (TCB) contains
thread-specific information

Stack pointer, PC, thread state
(running, ...), register values, a
pointer to PCB, ...

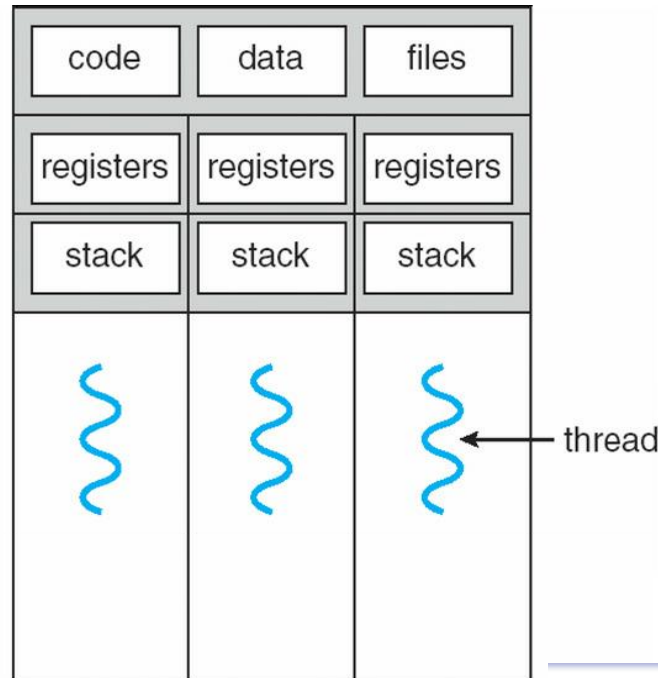




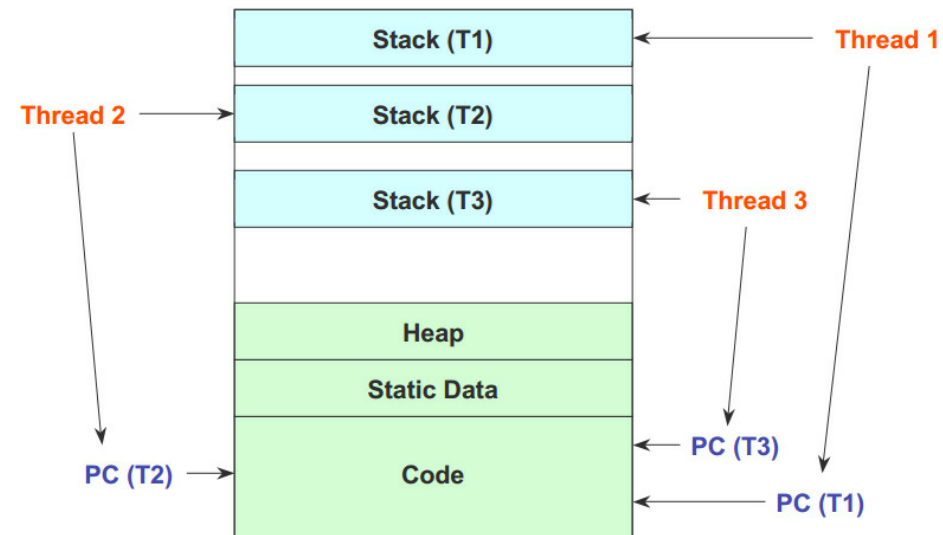
Single and Multithreaded Processes



single-threaded process



multithreaded process





Threads Benefits

A thread represents an abstract entity that executes a sequence of instructions

- It has its own set of CPU registers

- It has its own stack

- There is no thread-specific heap or data segment (unlike process)

Threads are lightweight

- Creating a thread more efficient than creating a process.

- Communication between threads easier than processes.

- Context switching between threads requires fewer CPU cycles and memory references than switching processes.

- Threads only track a subset of process state (share list of open files, pid, ...)

Examples:

- OS-supported: Windows' threads, Sun's LWP, POSIX threads





Context switch time for which entity is greater?

1. Process
2. Thread





Multithreads vs processor

P1	T1	C1+I1	$TT = C1 + I1 + C2$	Case1: Simple
			$TT = C1 + \text{Max}(I1, C2)$	Case2: Multithreaded process@uniprocessor
	T2	C2	$TT = \text{Max}(C1 + I1, C2)$	Case2: Multithreaded process@multiprocessors [Not work for Many to one Model]





How Can it Help?

Consider a Web server

Create a number of threads, and for each thread do

- get network message from client
- get URL data from disk
- send data over network

What did we gain?





Overlapping Requests (Concurrency)

Request 1
Thread 1

Request 2
Thread 2

get network message
(URL) from client
get URL data from disk
(disk access latency)

get network message
(URL) from client
~~get URL data from disk~~
(disk access latency)

send data over network

send data over network

- ◆ Total time is less than request 1 + request 2

Time





Threads vs. Processes

Threads

A thread has no separate data segment or heap

A thread cannot live on its own, it must live within a process

There can be more than one threads in a process, the first thread calls main & has the process's stack

If a thread dies, its stack is reclaimed

Inter-thread communication via memory.

Inexpensive creation and context switch

Processes

- ◆ A process has code/data/heap & other segments

- ◆ There must be at least one thread in a process

- ◆ Threads **within a process share** code/data/heap, share I/O, but each has its own stack & registers

- ◆ If a process dies, its resources are reclaimed & all threads die

- ◆ Inter-process communication via OS /data copying/message passing.


- ◆ Expensive creation and context switch





Threads' Life Cycle

Threads (just like processes) go through a sequence of *start*, *ready*, *running*, *waiting*, and *done* states





Threads have the same scheduling states as processes

1. True
2. False






Implementing Threads

POSIX Pthreads

User Level

- Threads package entirely in user space
- Kernel knows nothing about threads
- Fast to create and switch-scheduler as local procedure

Blocking System Call
CPU access




Windows NT Windows 2000

Kernel Level

- No runtime system
- Global Thread table, updated by kernel call
- Do not block process for syscall
- Thread switching: same or another process
- Not so fast as runtime system

Thread Recycling
Thread pool


Frequent thread operations
Need many System calls



Solaris

Hybrid

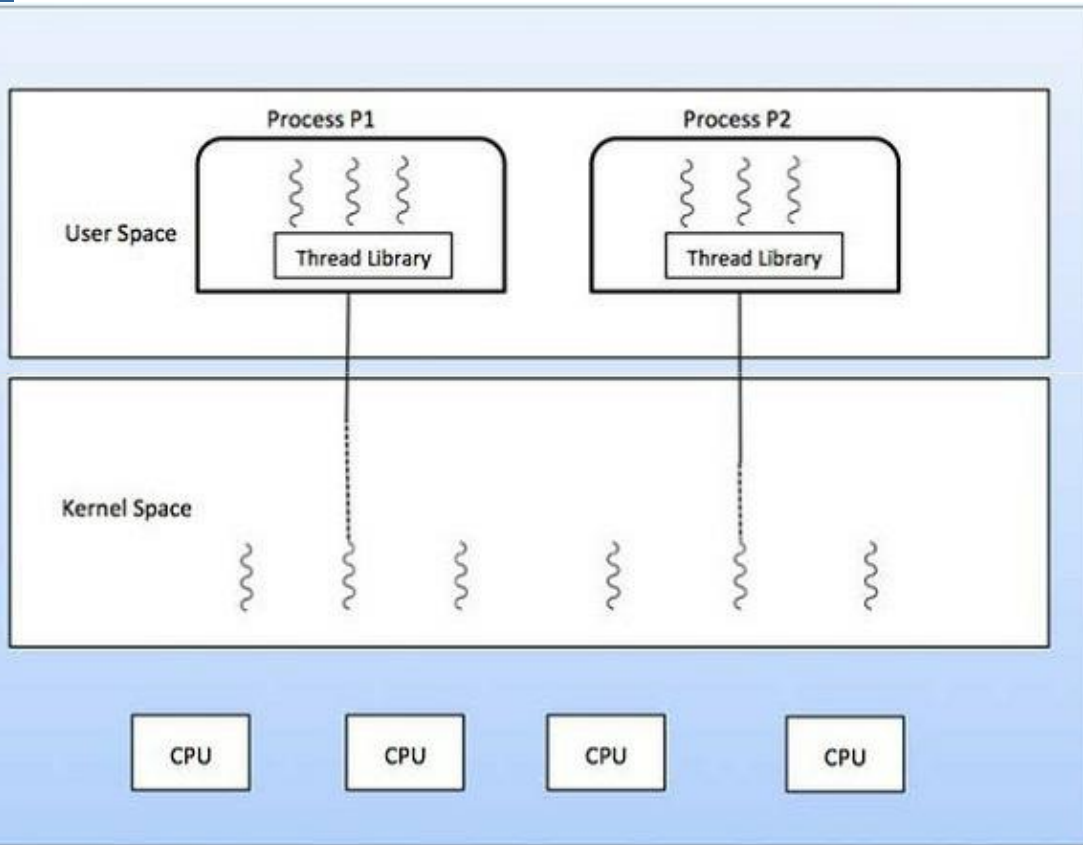
- Use kernel-level threads and then multiplex user-level threads onto same or all kernel threads.






Multithreading Models

Many to one Model



https://www.tutorialspoint.com/operating_system/os_multi_threading.htm

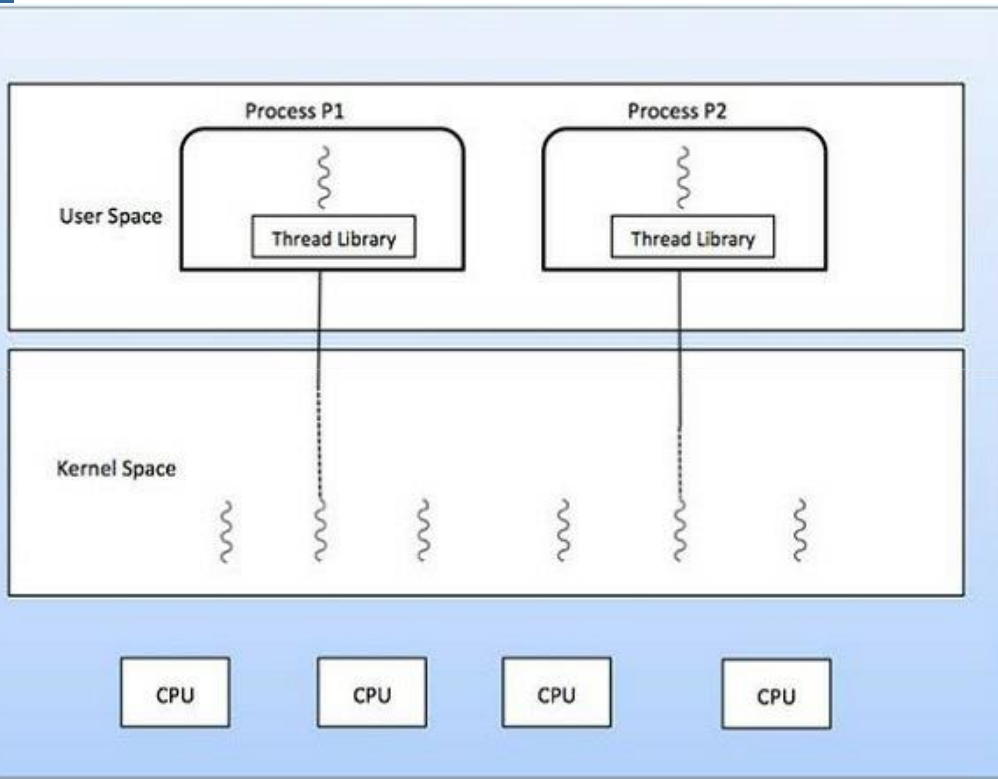
- Thread management is done in user space.
- Entire process is blocked if a thread makes a blocking system call.
- Not used in multiprocessor systems. No concurrency.
- Green threads – Solaris thread library.






Multithreading Models

One to one Model



https://www.tutorialspoint.com/operating_system/os_multi_threading.htm

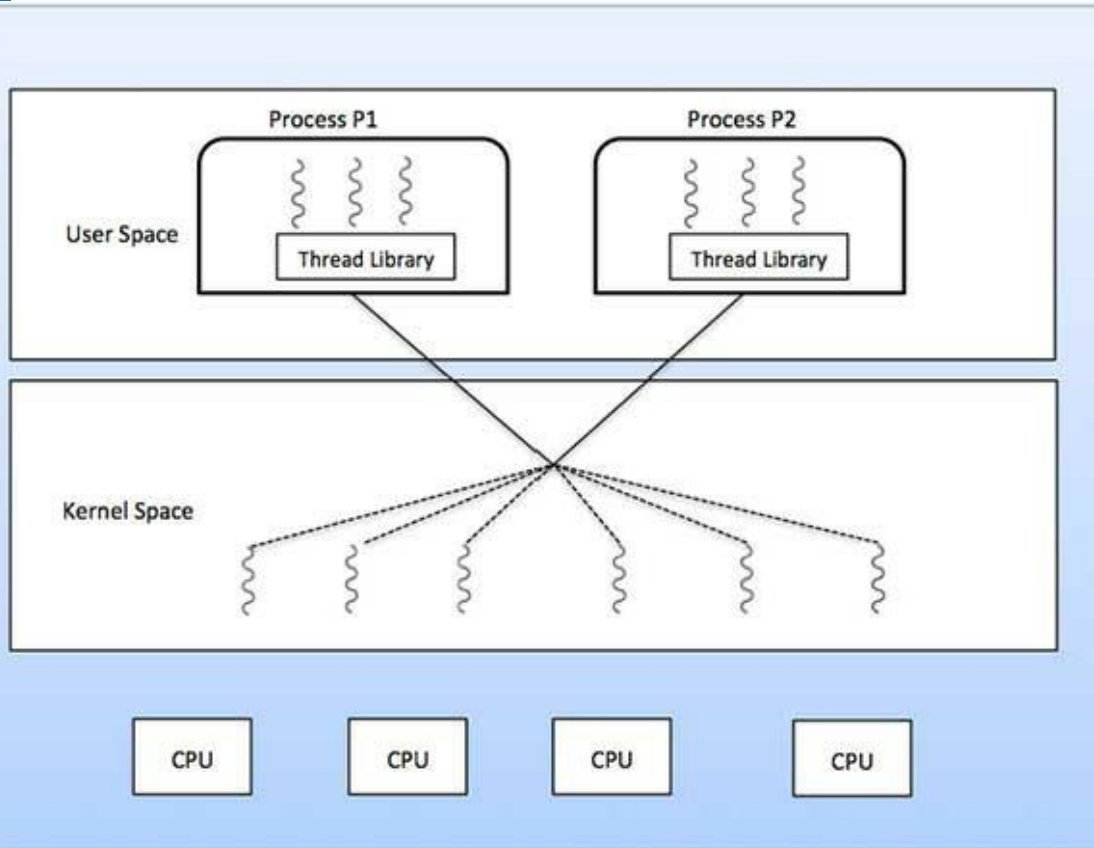
- Provide more concurrency than many to one.
 - another thread will run when a thread makes a blocking system call.
- Multiple threads to run in parallel on microprocessors.
- Drawback: creating a user thread, requires to create a kernel thread.
- Restriction has # User (or kernel) threads are supported by the system.
- Windows NT, Windows 2000






Multithreading Models

Many to Many Model



https://www.tutorialspoint.com/operating_system/os_multi_threading.htm

- Many users level threads to a **Greater** or equal number of kernel threads.
- Solve one-to-one, Many-to-one Model's restrictions.
- Users can create as many user threads as need
 - corresponding kernel threads can run parallel @ multiprocessor
- Solaris 2, IRIX, HP-UX



Course materials:
Galvin book 5.1-5.2,5.3.1

End of Lecture 6,7

